

A Model Of Predictive Governance

Jonas Hallgren (and potentially the king himself Loke Hagberg in the future??!!)

November 20, 2025

Abstract

Abstract: Social systems are complex, adaptive networks composed of overlapping hierarchies, markets, and communities. We do not claim to model this complexity in full. Instead, we apply a specific diagnostic lens: *Control Theory*. We view governance not as a static set of laws, but as a dynamic steering mechanism designed to maximize collective capability. We argue that traditional democratic feedback loops are often loosely coupled, leading to delayed error correction and low system viability. By viewing the polity as a "Macro-Agent" operating under the Free Energy Principle, we explore how mechanisms like *Predictive Governance* can tighten these loops. Crucially, we posit that for a system to be viable, it must integrate high-fidelity prediction signals into the existing "trust networks" of society via interpretable interfaces, thereby allowing the system to learn from simulated outcomes before real-world failures occur.

1 Introduction: The Goal of the System

What is the function of a governance system? If we look past ideological definitions regarding the "will of the people" and adopt a functionalist perspective, the goal of governance is to solve collective problems [11]. Whether the system is a small village managing an irrigation ditch or a nation-state managing an economy, the fundamental objective is to aggregate distributed knowledge to select actions that ensure the group's survival and flourishing [10].

However, as social groups scale, they face an entropy problem. Friedrich Hayek famously argued that the knowledge required to run a system is dispersed among all its members [8]. No central node possesses the requisite variety to model the whole. To manage this, societies build institutions—hierarchies, markets, and democracies—to compress this information.

Yet, the founder of Cybernetics, Norbert Wiener, noted that "effective behavior must be informed by some sort of feedback process" [15]. In modern governance, this feedback loop is often broken or severely delayed. We act (vote), but the feedback (outcome) arrives years later, buried in noise.

This paper argues that the crisis of modern governance is fundamentally a **Cybernetic Crisis**. The subsystems of society are poorly coupled. Information does not flow efficiently from the "Sense" organs to the "Decision" centers. To address this, we must look at the anatomy of the control loop itself.

2 The Control Loop: A Formal Description

To diagnose the problem, we must first define the ideal mechanism. Whether in a biological organism or a democratic state, effective steering requires a closed loop consisting of four distinct phases. This is not a reductionist metaphor, but a structural requirement for any goal-seeking agent.

Let us formally describe the process of a **Governing Agent** (the system):

1. **Sense (Observation):** The system gathers data y_t from the environment. In a society, this is done via decentralized networks (news, conversation) and price signals [9].
2. **Model (Prediction):** The system compares the current state to a preferred state (Goal). Crucially, based on the *Good Regulator Theorem* [4], the system must possess an internal model M of its environment. It generates a policy π predicting that this policy will lead to the highest reward (or lowest surprise) [5].
3. **Act (Implementation):** The system executes the policy. This is the vote, the law, or the resource allocation.
4. **Feedback (Reward/Update):** The environment changes in response to the action. The system observes the new state y_{t+1} .
 - If the outcome matches the prediction, the model M is reinforced.
 - If the outcome diverges, the system experiences a **Prediction Error** $\varepsilon = y_{t+1} - E[y_{t+1}|\pi]$.

The Goal: The system wants to update its model M such that future actions maximize reward. This is the process of *Learning*. If the loop is too slow, the model M drifts away from reality, and the system becomes non-viable.




Figure 1: **The Governance Control Loop.** For a system to be viable, the feedback signal must be tightly coupled to the sensing mechanism. In many current systems, this link is broken by latency (time delay) and noise, preventing the system from learning.

3 Structural Diagnosis: The Coupling Problem

If the theoretical loop is clear, where does reality fail? We turn to Stafford Beer’s *Viable System Model* (VSM) for a diagnostic lens [2]. Beer argued that for a system to survive, it must manage the interplay between its internal operations and the external environment through specific functional subsystems.

In modern democracies, we identify three critical ”Coupling Failures” where the connection between these subsystems breaks down:

3.1 Latency (The Time Problem)

The feedback loop often takes years (e.g., a 4-year election cycle). A policy is enacted, but the ”Reward” signal arrives so late that the system cannot causally link the specific action to the specific outcome. In Active Inference terms, the *temporal depth* of the model is insufficient for the complexity of the environment [12].

3.2 Variety (The Complexity Problem)

As per Ashby’s Law of Requisite Variety [1], the governance system must have as much variety (possible states) as the system it regulates. Current systems often attempt to

regulate complex, high-variety economies with low-variety binary voting mechanisms. This leads to "attenuation" where critical nuances are lost.

3.3 Information Flow (The Network Problem)

Information flows primarily through *Social Trust Networks*. As modeled by Network Science, these "Small World" networks are efficient for cohesion [14]. However, they act as filters. If a high-fidelity error signal ("This policy is failing") originates from outside a trust cluster, the cluster often blocks the signal to preserve internal stability [13].

We posit that the "Broken Link" in Figure 1 is a failure of **Prediction**. The system acts on *Intent* ("I want a good outcome") rather than a rigorous *Model of the Future* ("This action causes that outcome").

4 Closing the Loop: Predictive Governance

To fix the coupling, we cannot simply "remove" the humans or the trust networks; they are the substrate of society. Instead, we must *augment* the loop.

We introduce **Predictive Governance** (and specifically concepts like Predictive Liquid Democracy [6] and Futarchy [7]) not as ideological replacements for democracy, but as cybernetic mechanisms to tighten the feedback loop.

By utilizing prediction markets or forecasting ensembles, the system generates a **Simulated Feedback** signal.

- **Traditional Loop:** Act → Wait 4 Years → Real World Feedback.
- **Predictive Loop:** Proposal → Simulation/Market → Immediate Feedback Signal → Act.

This creates a "Feed-forward" mechanism [3]. It allows the system to "learn" from virtual mistakes before committing real resources, effectively maximizing the expected reward and minimizing the free energy of the collective agent.

5 Formalizing the Polity as an Active Inference Agent

To engineer a better control loop, we must move beyond metaphors and define the system mathematically. We model the polity (the collective decision-making body) as a "Macro-

Agent” operating under the *Free Energy Principle* [5].

5.1 The Generative Model of Governance

The polity maintains a Generative Model \mathcal{M} that encapsulates its beliefs about how policies lead to outcomes.

$$\mathcal{M} = \{P(o, s, \pi)\}$$

Where:

- o are **Observations** (economic indicators, social well-being).
- s are **Hidden States** (the true condition of the world).
- π are **Policies** (laws, budget allocations).

5.2 The Objective: Minimizing Expected Free Energy

The goal of the governance system is to select a policy π that minimizes *Expected Free Energy* $G(\pi)$. In Active Inference, $G(\pi)$ is composed of two distinct terms [12]:

$$G(\pi) = \underbrace{D_{KL}[Q(o|\pi)||P(o)]}_{\text{Risk (Pragmatic)}} + \underbrace{\mathbb{E}_Q[\ln Q(s|o, \pi) - \ln Q(s|\pi)]}_{\text{Ambiguity (Epistemic)}} \quad (1)$$

1. **Risk (Pragmatic Value):** The divergence between predicted outcomes and *preferred* outcomes. (e.g., ”We prefer low inflation. Does Policy A predict low inflation?”).
2. **Ambiguity (Epistemic Value):** The drive to resolve uncertainty. (e.g., ”We don’t know what Policy A does. We need to find out.”).

The Failure of the 4-Year Controller: Current democratic systems utilize a ”4-Year Controller.” The temporal depth of the loop is $\tau = 4$ years. This means the system only evaluates $G(\pi)$ once per cycle. Between cycles, the system is ”blind,” accumulating variational free energy (surprise) as the environment changes but the policy remains static [3].

6 Engineering the Fast Loop: The Mechanics of Predictive Liquid Democracy

Having established the need for a control loop with higher temporal resolution, we now detail the specific architecture of Predictive Liquid Democracy (PLD). We frame PLD not as a political ideology, but as a **Cybernetic Circuit** designed to minimize the accumulation of variational free energy over time.

6.1 The Problem of Temporal Depth

In Active Inference, an agent plans behavior based on a "Temporal Horizon" or depth τ [12]. The agent projects its model into the future to estimate Expected Free Energy $G(\pi)$.

Standard Democracy (The 4-Year Sample): In a traditional representative system, the sampling rate of the collective agent is extremely low ($\approx 0.7 \times 10^{-8}$ Hz, or once every 4 years).

- t_0 : Policy enacted.
- $t_0 \rightarrow t_4$: The environment changes. The policy drifts from optimality. The system accumulates error (Surprise).
- t_4 : Feedback is received (Election). Correction is applied.

This creates a "Sawtooth" error function, where divergence maximizes just before the correction.

PLD (The Continuous Controller): PLD introduces a continuous-time estimator. By using prediction markets, the system samples the expected outcome $E[o|\pi]$ continuously. This effectively reduces the feedback delay $\Delta t \rightarrow 0$.

6.2 Component A: The Virtual Feedback Loop (System 4)

To achieve continuous correction, we cannot wait for reality. We must simulate it. The Prediction Market serves as the **Generative Model Simulator**.

Let π be a proposed policy. Let O_{goal} be the preferred outcome (e.g., GDP Growth i 2%). The market generates a price signal P_m representing the conditional probability:

$$P_m(\pi) \approx P(O_{goal}|\pi) \tag{2}$$




Figure 2: **The Accumulation of Error.** Standard governance (Red) corrects error only at discrete intervals (elections), leading to high integrated free energy. PLD (Blue) uses real-time predictive feedback to micro-correct policies and delegations, keeping the system continuously aligned with its goals.

This signal acts as **Virtual Sensory Data**. In the VSM, this constitutes System 4 (Intelligence). It forwards a "Future Error Signal" ε_{future} to the control center before the error actually occurs.

6.3 Component B: The Dynamic Actuator (System 3)

Knowing the error is insufficient; the system must act. In PLD, the "Actuator" is the **Delegation Graph \mathcal{G}** .

Let $W_{ij}(t)$ be the weight of the vote delegated from voter i to delegate j at time t . In a static democracy, W is frozen for the term. In PLD, W is dynamic. The update rule for the delegation graph can be formalized as a function of the Virtual Error Signal ε_{future} :

$$W_{ij}(t + 1) = W_{ij}(t) - \alpha \cdot \text{AlignmentError}(j, \varepsilon_{future}) \quad (3)$$

Where:

- α is the learning rate (agility of the electorate).
- AlignmentError measures the specific delegate's support for a policy predicted to fail.

If Delegate j votes for a policy that the Market predicts will fail, AlignmentError spikes. The system (or the user via dashboard) automatically reallocates W_{ij} to a delegate with lower error.




Figure 3: **The PLD Circuit Diagram.** This diagram illustrates the "Short Loop." 1. A policy is proposed. 2. The Market simulates the outcome. 3. The Comparator checks if the predicted outcome matches the Goal. 4. If there is a gap (Error), the User Dashboard signals a change in delegation. 5. The Delegation Graph updates *before* the policy is fully implemented in the real world.

6.4 Deep Coupling: From Prediction to Trust

The innovation of this architecture is how it handles the **Epistemic vs. Pragmatic** trade-off.

- **Markets minimize Epistemic Ambiguity:** They answer "What will happen?"
- **Delegation minimizes Complexity:** It answers "Who should I follow?"

In a standard system, voters stick with bad delegates because evaluating them is too costly (High Complexity). By providing a clear, low-latency error signal ε (via the Market), PLD reduces the cognitive cost of switching delegates.

The system essentially says: “*You don’t need to understand the policy (High Cost). You only need to see that your Delegate is betting against the Market (Low Cost).*” This allows the Graph W to converge toward optimal configuration faster than the rate of environmental change, satisfying Ashby’s Law of Requisite Variety.

7 The Network Topology of Trust: Integrating Hybrid Signals

We have established that PLD shortens the temporal loop. However, governance is not just a temporal problem; it is a **Spatial Problem** defined by the topology of the social graph.

In this section, we explore how the ”Adversarial” nature of prediction markets interacts with the ”Cooperative” nature of social trust, and how this hybrid architecture affects broader democratic models beyond PLD.

7.1 The Graph of Priors: Why Deliberation Matters

In Active Inference, an agent’s worldview is constrained by its **Markov Blanket**—it can only know what its neighbors tell it.

$$\text{Social Graph } \mathcal{G} = (V, E_{social})$$

Edges in E_{social} represent **Trust Channels**. These are efficient for transmitting values and maintaining cohesion (System 5 dynamics). However, Network Science shows that these graphs are highly clustered (homophilic) [14]. This leads to *Epistemic Closure*: if a cluster holds a false belief, the strong trust links reinforce that error, blocking external correction [13].

The Role of Deliberation: Deliberative Democracy functions as a mechanism to create temporary ”weak ties” between clusters, allowing for the cross-pollination of priors. However, without a ground-truth signal, deliberation can sometimes reinforce group polarization.

7.2 The Market as a Broadcast Channel

The Prediction Market functions differently. It is not a peer-to-peer link; it is a **Global Broadcast Channel**.

$$\text{Signal } S_{market} \rightarrow \forall v \in V$$

Unlike social links, the market is **Adversarial**. Traders compete (zero-sum) to correct errors. This competition is what generates the *High Precision* (ω_{market}) of the signal. It is an "Information Pump" that pulls private information from the edges of the network and broadcasts it to the center [8].

7.3 The Hybrid Update: Rewiring the Graph

The core challenge of the "Hybrid System" is: How do we inject this global, adversarial signal into local, cooperative clusters without destroying them?

We model this as a **Precision-Weighted Network Update**. The Agent i receives two signals regarding a Policy π :

1. **Social Signal** (y_{soc}): Weighted by Trust ω_{soc} (High for friends).
2. **Market Signal** (y_{mkt}): Weighted by Accuracy ω_{mkt} (Initially low for average users).

The Bridge Node Mechanism: Most voters will ignore y_{mkt} due to complexity. However, the graph contains "Bridge Nodes" (or Opinion Leaders)—individuals with high cognitive bandwidth who pay attention to the Market (System 4).


When the Market signals an error (ε), these Bridge Nodes update their priors first. Because they are trusted by their local clusters, they propagate this updated prior down the social edges.

Result: The Prediction Market does not need to convince every voter. It only needs to convince the Bridge Nodes. The "Trust Network" then does the work of propagating the truth to the edges.

8 Conclusion: Toward Generalized Active Governance

We began this inquiry with a simple question: How can we improve the decision-making capability of a democratic system? By applying the lens of **Control Theory** and **Active Inference**, we have arrived at a generalized framework for governance.

We conclude with three implications for the future of political economy:



The Bridge Mechanism

Market signals are absorbed by "Bridge Nodes" (Delegates) who verify the data and propagate it via Social Trust.

Figure 4: **Topology of Information Flow.** Voters (Grey) are shielded from the complexity of the Market (Red) by their cognitive bounds. "Bridge Nodes" (Blue)—who effectively act as System 3 managers—absorb the high-fidelity signal from the Market and translate it into "Trust" signals for their local clusters. This combines the accuracy of the market with the cohesion of the social graph.

8.1 1. The Cybernetic State is a Learning State

Viability is not a static property; it is a dynamic capability. A state is only viable if it can minimize its variational free energy (surprise) over time. This requires feedback loops that are **fast** (high temporal resolution) and **accurate** (high signal fidelity). Tools like PLD are best understood not as ideological disruptions, but as necessary upgrades to the "sensory circuitry" of the state.

8.2 2. Interpretability is a Thermodynamic Constraint

We cannot simply optimize governance using black-box AI or incomprehensible markets. The "Human-in-the-Loop" is not just an ethical requirement; it is a physical one. For a social system to maintain low entropy, its constituent agents must be able to model the system's behavior. **Interpretability** reduces the metabolic cost of trust. Therefore, any predictive governance system must prioritize the "Glass Box" interface—providing narratives, not just numbers.

8.3 3. Hybirdity is the End State

The debate between "Markets" vs. "Democracy" vs. "Technocracy" is a category error. A viable system requires all three:

- **Markets (Adversarial Sensors)** to generate high-precision information (System 4).
- **Networks (Cooperative Trust)** to maintain cohesion and value alignment (System 5).
- **Hierarchy (Bridge Nodes)** to translate between the two (System 3).

The future of governance lies in designing architectures that allow these distinct mechanisms to couple tightly, creating a "Macro-Agent" that is smarter than its parts and more humane than its algorithms.

References

- [1] W. Ross Ashby. *An Introduction to Cybernetics*. Chapman & Hall, 1956.
- [2] Stafford Beer. *Brain of the Firm*. Allen Lane, 1972.
- [3] Stafford Beer. *Designing Freedom*. House of Anansi Press, 1974.
- [4] Roger C Conant and W Ross Ashby. Every good regulator of a system must be a model of that system. *International Journal of Systems Science*, 1(2):89–97, 1970.
- [5] Karl Friston. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010.
- [6] Loke Hagberg and Samuel Färdow Kazen. Predictive liquid democracy. *Preprint*, 2024.
- [7] Robin Hanson. Shall we vote on values, but bet on beliefs? *Journal of Political Philosophy*, 21(2):151–178, 2013.
- [8] Friedrich A Hayek. The use of knowledge in society. *The American Economic Review*, 35(4):519–530, 1945.
- [9] Thomas W Malone. *Superminds: The Surprising Power of People and Computers Thinking Together*. Little, Brown, 2018.
- [10] Josiah Ober. *Democracy and Knowledge: Innovation and Learning in Classical Athens*. Princeton University Press, 2008.
- [11] Elinor Ostrom. *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press, 1990.
- [12] Thomas Parr, Giovanni Pezzulo, and Karl J Friston. *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. MIT Press, 2022.
- [13] Alex Pentland. *Social Physics: How Good Ideas Spread-The Lessons from a New Science*. Penguin, 2014.
- [14] Duncan J Watts and Steven H Strogatz. Collective dynamics of ‘small-world’ networks. *nature*, 393(6684):440–442, 1998.
- [15] Norbert Wiener. *Cybernetics: Or Control and Communication in the Animal and the Machine*. MIT Press, 1948.